

CS110 Lecture 09: Threads

Principles of Computer Systems

Winter 2020

Stanford University

Computer Science Department

Instructors: Chris Gregg and
Nick Troccoli

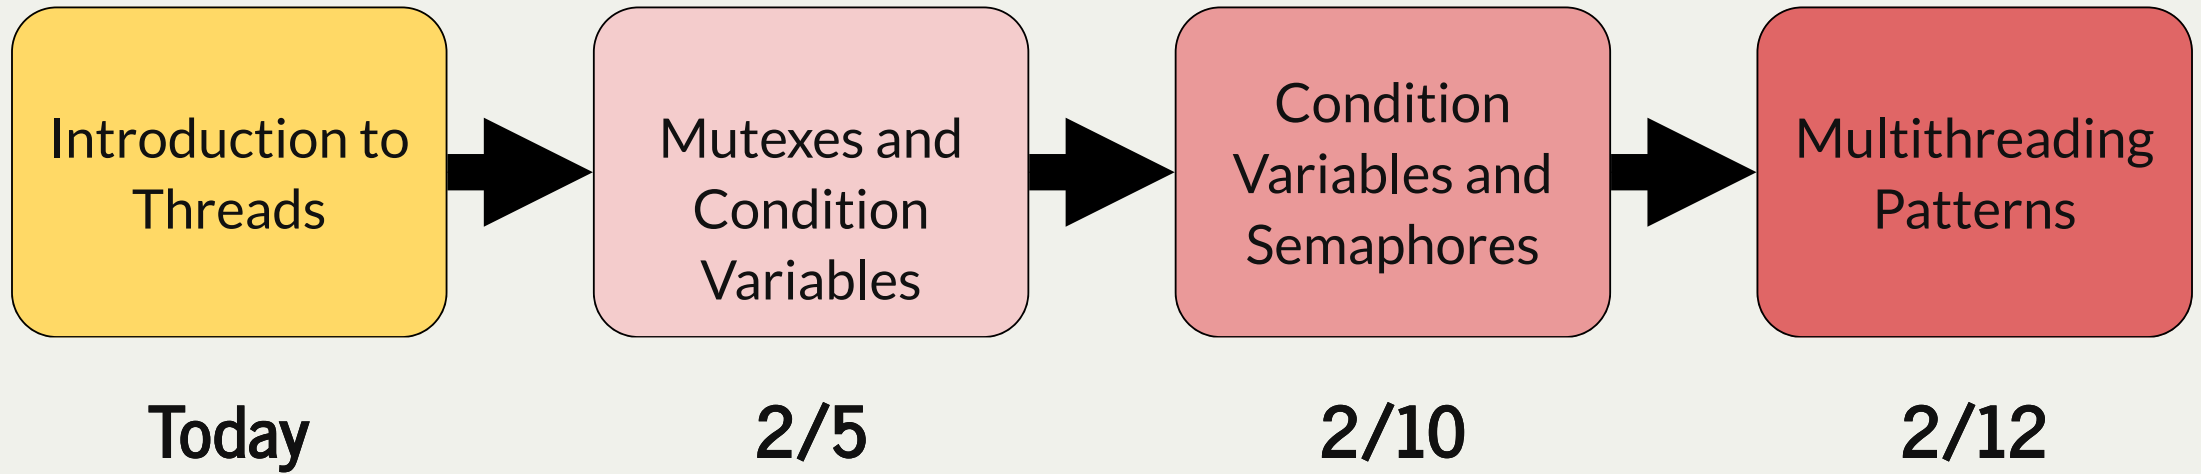


[PDF of this presentation](#)

CS110 Topic 3: How can we have concurrency within a single process?



Learning About Processes



Today's Learning Goals

- Learn about how threads allow for concurrency within a single process
- Understand the differences between threads and processes
- Discover some of the pitfalls of threads sharing the same virtual address space



Plan For Today

- More Practice: Race Conditions
- **Break:** Announcements
- Threads



Practice Problem 1

Consider this program and its execution. Assume that all processes run to completion, all system and `printf` calls succeed, and that all calls to `printf` are atomic. Assume nothing about scheduling or time slice durations.

```
1 static void bat(int unused) {  
2     printf("pirate\n");  
3     exit(0);  
4 }  
5  
6 int main(int argc, char *argv[]) {  
7     signal(SIGUSR1, bat);  
8     pid_t pid = fork();  
9     if (pid == 0) {  
10         printf("ghost\n");  
11         return 0;  
12     }  
13     kill(pid, SIGUSR1);  
14     printf("ninja\n"); return 0;  
15 }
```

- For each of the five output orders, Place a **yes** if the text represents a possible output, and place a **no** otherwise.
 - ghost ninja pirate
 - pirate ninja
 - ninja ghost
 - ninja pirate ninja
 - ninja pirate ghost



Practice Problem 1

Consider this program and its execution. Assume that all processes run to completion, all system and `printf` calls succeed, and that all calls to `printf` are atomic. Assume nothing about scheduling or time slice durations.

```
1 static void bat(int unused) {  
2     printf("pirate\n");  
3     exit(0);  
4 }  
5  
6 int main(int argc, char *argv[]) {  
7     signal(SIGUSR1, bat);  
8     pid_t pid = fork();  
9     if (pid == 0) {  
10         printf("ghost\n");  
11         return 0;  
12     }  
13     kill(pid, SIGUSR1);  
14     printf("ninja\n"); return 0;  
15 }
```

- For each of the five output orders, Place a **yes** if the text represents a possible output, and place a **no** otherwise.
 - ghost ninja pirate. **yes**
 - pirate ninja. **yes**
 - ninja ghost. **no**
 - ninja pirate ninja. **no**
 - ninja pirate ghost. **no**



Practice Problem 2

Consider this program and its execution. Assume that all processes run to completion, all system and `printf` calls succeed, and that all calls to `printf` are atomic. Assume nothing about scheduling or time slice durations.

```
1 int main(int argc, char *argv[]) {
2     pid_t pid;
3     int counter = 0;
4     while (counter < 2) {
5         pid = fork();
6         if (pid > 0) break;
7         counter++;
8         printf("%d", counter);
9     }
10    if (counter > 0) printf("%d", counter);
11    if (pid > 0) {
12        waitpid(pid, NULL, 0);
13        counter += 5;
14        printf("%d", counter);
15    }
16    return 0;
17 }
```

- List all possible outputs
 - How many processes are created in total?
 - What number must be printed last?
 - What number must be printed first?
 - What number must be printed second-last?



Practice Problem 2

Consider this program and its execution. Assume that all processes run to completion, all system and `printf` calls succeed, and that all calls to `printf` are atomic. Assume nothing about scheduling or time slice durations.

```
1 int main(int argc, char *argv[]) {
2     pid_t pid;
3     int counter = 0;
4     while (counter < 2) {
5         pid = fork();
6         if (pid > 0) break;
7         counter++;
8         printf("%d", counter);
9     }
10    if (counter > 0) printf("%d", counter);
11    if (pid > 0) {
12        waitpid(pid, NULL, 0);
13        counter += 5;
14        printf("%d", counter);
15    }
16    return 0;
17 }
```

- List all possible outputs
 - How many processes are created in total? 3
 - What number must be printed last?
 - What number must be printed first?
 - What number must be printed second-last?



Practice Problem 2

Consider this program and its execution. Assume that all processes run to completion, all system and `printf` calls succeed, and that all calls to `printf` are atomic. Assume nothing about scheduling or time slice durations.

```
1 int main(int argc, char *argv[]) {
2     pid_t pid;
3     int counter = 0;
4     while (counter < 2) {
5         pid = fork();
6         if (pid > 0) break;
7         counter++;
8         printf("%d", counter);
9     }
10    if (counter > 0) printf("%d", counter);
11    if (pid > 0) {
12        waitpid(pid, NULL, 0);
13        counter += 5;
14        printf("%d", counter);
15    }
16    return 0;
17 }
```

- List all possible outputs
 - What number must be printed last?
 - The last output is 5 (the grandparent that starts with counter 0 exits last, because it waits on its child)
 - What number must be printed first?
 - What number must be printed second-last?



Practice Problem 2

Consider this program and its execution. Assume that all processes run to completion, all system and `printf` calls succeed, and that all calls to `printf` are atomic. Assume nothing about scheduling or time slice durations.

```
1 int main(int argc, char *argv[]) {
2     pid_t pid;
3     int counter = 0;
4     while (counter < 2) {
5         pid = fork();
6         if (pid > 0) break;
7         counter++;
8         printf("%d", counter);
9     }
10    if (counter > 0) printf("%d", counter);
11    if (pid > 0) {
12        waitpid(pid, NULL, 0);
13        counter += 5;
14        printf("%d", counter);
15    }
16    return 0;
17 }
```

- List all possible outputs
 - What number must be printed last?
 - The last output is 5 (the grandparent that starts with counter 0 exits last, because it waits on its child)
 - What number must be printed first?
 - The parent that starts with counter 1 outputs first
 - What number must be printed second-last?



Practice Problem 2

Consider this program and its execution. Assume that all processes run to completion, all system and `printf` calls succeed, and that all calls to `printf` are atomic. Assume nothing about scheduling or time slice durations.

```
1 int main(int argc, char *argv[]) {
2     pid_t pid;
3     int counter = 0;
4     while (counter < 2) {
5         pid = fork();
6         if (pid > 0) break;
7         counter++;
8         printf("%d", counter);
9     }
10    if (counter > 0) printf("%d", counter);
11    if (pid > 0) {
12        waitpid(pid, NULL, 0);
13        counter += 5;
14        printf("%d", counter);
15    }
16    return 0;
17 }
```

- List all possible outputs
 - What number must be printed last?
 - The last output is 5 (the grandparent that starts with counter 0 exits last, because it waits on its child)
 - What number must be printed first?
 - The parent that starts with counter 1 outputs first
 - What number must be printed second-last?
 - The parent that starts with counter 1 exits second to last, after waiting for its child; the second-to-last output is 6

Practice Problem 2

Consider this program and its execution. Assume that all processes run to completion, all system and `printf` calls succeed, and that all calls to `printf` are atomic. Assume nothing about scheduling or time slice durations.

```
1 int main(int argc, char *argv[]) {
2     pid_t pid;
3     int counter = 0;
4     while (counter < 2) {
5         pid = fork();
6         if (pid > 0) break;
7         counter++;
8         printf("%d", counter);
9     }
10    if (counter > 0) printf("%d", counter);
11    if (pid > 0) {
12        waitpid(pid, NULL, 0);
13        counter += 5;
14        printf("%d", counter);
15    }
16    return 0;
17 }
```

- List all possible outputs
 - Possible Output 1: 112265
 - Possible Output 2: 121265
 - Possible Output 3: 122165



Practice Problem 2

Consider this program and its execution. Assume that all processes run to completion, all system and `printf` calls succeed, and that all calls to `printf` are atomic. Assume nothing about scheduling or time slice durations.

```
1 int main(int argc, char *argv[]) {
2     pid_t pid;
3     int counter = 0;
4     while (counter < 2) {
5         pid = fork();
6         if (pid > 0) break;
7         counter++;
8         printf("%d", counter);
9     }
10    if (counter > 0) printf("%d", counter);
11    if (pid > 0) {
12        waitpid(pid, NULL, 0);
13        counter += 5;
14        printf("%d", counter);
15    }
16    return 0;
17 }
```

- List all possible outputs
 - Possible Output 1: 112265
 - Possible Output 2: 121265
 - Possible Output 3: 122165
- The second parent (1) output and the two child (2) outputs are up to the scheduler



Practice Problem 2

Consider this program and its execution. Assume that all processes run to completion, all system and `printf` calls succeed, and that all calls to `printf` are atomic. Assume nothing about scheduling or time slice durations.

```
1 int main(int argc, char *argv[]) {
2     pid_t pid;
3     int counter = 0;
4     while (counter < 2) {
5         pid = fork();
6         if (pid > 0) break;
7         counter++;
8         printf("%d", counter);
9     }
10    if (counter > 0) printf("%d", counter);
11    if (pid > 0) {
12        waitpid(pid, NULL, 0);
13        counter += 5;
14        printf("%d", counter);
15    }
16    return 0;
17 }
```

- List all possible outputs
 - Possible Output 1: 112265
 - Possible Output 2: 121265
 - Possible Output 3: 122165
- If the `>` of the `counter > 0` test is changed to a `>=`, then `counter` values of zeroes would be included in each possible output. How many different outputs are now possible? (No need to list the outputs—just present the number.)



Practice Problem 2

Consider this program and its execution. Assume that all processes run to completion, all system and `printf` calls succeed, and that all calls to `printf` are atomic. Assume nothing about scheduling or time slice durations.

```
1 int main(int argc, char *argv[]) {
2     pid_t pid;
3     int counter = 0;
4     while (counter < 2) {
5         pid = fork();
6         if (pid > 0) break;
7         counter++;
8         printf("%d", counter);
9     }
10    if (counter > 0) printf("%d", counter);
11    if (pid > 0) {
12        waitpid(pid, NULL, 0);
13        counter += 5;
14        printf("%d", counter);
15    }
16    return 0;
17 }
```

- List all possible outputs
 - Possible Output 1: 112265
 - Possible Output 2: 121265
 - Possible Output 3: 122165
- If the `>` of the `counter > 0` test is changed to a `>=`, then `counter` values of zeroes would be included in each possible output. How many different outputs are now possible? (No need to list the outputs—just present the number.) **18 - 6x the original**



Practice Problem 3

Consider the following program. Assume that each call to **printf** flushes its output to the console in full, and further assume that none of the system calls fail in any unpredictable way.

```
1 static pid_t pid; // necessarily global so handler1 has access
2 static int counter = 0;
3 static void handler1(int unused) {
4     counter++;
5     printf("counter = %d\n", counter);
6     kill(pid, SIGUSR1);
7 }
8 static void handler2(int unused) {
9     counter += 10;
10    printf("counter = %d\n", counter);
11    exit(0);
12 }
13 int main(int argc, char *argv[]) {
14     signal(SIGUSR1, handler1);
15     if ((pid = fork()) == 0) {
16         signal(SIGUSR1, handler2);
17         kill(getppid(), SIGUSR1);
18         while (true) {}
19     }
20     if (waitpid(-1, NULL, 0) > 0) {
21         counter += 1000;
22         printf("counter = %d\n", counter);
23     }
24     return 0;
25 }
```

- What is the output of the program?



Practice Problem 3

Consider the following program. Assume that each call to **printf** flushes its output to the console in full, and further assume that none of the system calls fail in any unpredictable way.

```
1 static pid_t pid; // necessarily global so handler1 has access
2 static int counter = 0;
3 static void handler1(int unused) {
4     counter++;
5     printf("counter = %d\n", counter);
6     kill(pid, SIGUSR1);
7 }
8 static void handler2(int unused) {
9     counter += 10;
10    printf("counter = %d\n", counter);
11    exit(0);
12 }
13 int main(int argc, char *argv[]) {
14     signal(SIGUSR1, handler1);
15     if ((pid = fork()) == 0) {
16         signal(SIGUSR1, handler2);
17         kill(getppid(), SIGUSR1);
18         while (true) {}
19     }
20     if (waitpid(-1, NULL, 0) > 0) {
21         counter += 1000;
22         printf("counter = %d\n", counter);
23     }
24     return 0;
25 }
```

- What is the output of the program?

counter = 1

counter = 10

counter = 1001



Practice Problem 3

Consider the following program. Assume that each call to **printf** flushes its output to the console in full, and further assume that none of the system calls fail in any unpredictable way.

```
1 static pid_t pid; // necessarily global so handler1 has access
2 static int counter = 0;
3 static void handler1(int unused) {
4     counter++;
5     printf("counter = %d\n", counter);
6     kill(pid, SIGUSR1);
7 }
8 static void handler2(int unused) {
9     counter += 10;
10    printf("counter = %d\n", counter);
11    exit(0);
12 }
13 int main(int argc, char *argv[]) {
14     signal(SIGUSR1, handler1);
15     if ((pid = fork()) == 0) {
16         signal(SIGUSR1, handler2);
17         kill(getppid(), SIGUSR1);
18         while (true) {}
19     }
20     if (waitpid(-1, NULL, 0) > 0) {
21         counter += 1000;
22         printf("counter = %d\n", counter);
23     }
24     return 0;
25 }
```

- What are the two potential outputs of the above program if the **while (true)** loop is completely eliminated?



Practice Problem 3

Consider the following program. Assume that each call to **printf** flushes its output to the console in full, and further assume that none of the system calls fail in any unpredictable way.

```
1 static pid_t pid; // necessarily global so handler1 has access
2 static int counter = 0;
3 static void handler1(int unused) {
4     counter++;
5     printf("counter = %d\n", counter);
6     kill(pid, SIGUSR1);
7 }
8 static void handler2(int unused) {
9     counter += 10;
10    printf("counter = %d\n", counter);
11    exit(0);
12 }
13 int main(int argc, char *argv[]) {
14     signal(SIGUSR1, handler1);
15     if ((pid = fork()) == 0) {
16         signal(SIGUSR1, handler2);
17         kill(getppid(), SIGUSR1);
18         while (true) {}
19     }
20     if (waitpid(-1, NULL, 0) > 0) {
21         counter += 1000;
22         printf("counter = %d\n", counter);
23     }
24     return 0;
25 }
```

- What are the two potential outputs of the above program if the **while (true)** loop is completely eliminated?
 - Previous output still possible via scheduling
 - Now, though, the child process could complete and exit normally before the parent process— via its **handler1** function— has the opportunity to signal the child. That would mean **handler2** wouldn't even execute.
 - So, another possible output would be:

counter = 1
counter = 1001



Practice Problem 3

Consider the following program. Assume that each call to **printf** flushes its output to the console in full, and further assume that none of the system calls fail in any unpredictable way.

```
1 static pid_t pid; // necessarily global so handler1 has access
2 static int counter = 0;
3 static void handler1(int unused) {
4     counter++;
5     printf("counter = %d\n", counter);
6     kill(pid, SIGUSR1);
7 }
8 static void handler2(int unused) {
9     counter += 10;
10    printf("counter = %d\n", counter);
11    exit(0);
12 }
13 int main(int argc, char *argv[]) {
14     signal(SIGUSR1, handler1);
15     if ((pid = fork()) == 0) {
16         signal(SIGUSR1, handler2);
17         kill(getppid(), SIGUSR1);
18         while (true) {}
19     }
20     if (waitpid(-1, NULL, 0) > 0) {
21         counter += 1000;
22         printf("counter = %d\n", counter);
23     }
24     return 0;
25 }
```

- Now further assume the call to **exit(0)** has also been removed from the **handler2** function. Are there any other potential program outputs? If not, explain why. If so, what are they?
 - No other potential outputs, because:
 - **counter = 1** is still printed exactly once, just in the parent, before the parent fires a **SIGUSR1** signal at the child
 - ...



Practice Problem 3

Consider the following program. Assume that each call to **printf** flushes its output to the console in full, and further assume that none of the system calls fail in any unpredictable way.

```
1 static pid_t pid; // necessarily global so handler1 has access
2 static int counter = 0;
3 static void handler1(int unused) {
4     counter++;
5     printf("counter = %d\n", counter);
6     kill(pid, SIGUSR1);
7 }
8 static void handler2(int unused) {
9     counter += 10;
10    printf("counter = %d\n", counter);
11    exit(0);
12 }
13 int main(int argc, char *argv[]) {
14     signal(SIGUSR1, handler1);
15     if ((pid = fork()) == 0) {
16         signal(SIGUSR1, handler2);
17         kill(getppid(), SIGUSR1);
18         while (true) {}
19     }
20     if (waitpid(-1, NULL, 0) > 0) {
21         counter += 1000;
22         printf("counter = %d\n", counter);
23     }
24     return 0;
25 }
```

- Now further assume the call to **exit(0)** has also been removed from the **handler2** function. Are there any other potential program outputs? If not, explain why. If so, what are they?
 - No other potential outputs, because:
 - ...
 - **counter = 10** is potentially printed if the child is still running at the time the parent fires that **SIGUSR1** signal at it. The **10** can only appear after the **1**, and if it appears, it must appear before the **1001**.
 - ...

Practice Problem 3

Consider the following program. Assume that each call to **printf** flushes its output to the console in full, and further assume that none of the system calls fail in any unpredictable way.

```
1 static pid_t pid; // necessarily global so handler1 has access
2 static int counter = 0;
3 static void handler1(int unused) {
4     counter++;
5     printf("counter = %d\n", counter);
6     kill(pid, SIGUSR1);
7 }
8 static void handler2(int unused) {
9     counter += 10;
10    printf("counter = %d\n", counter);
11    exit(0);
12 }
13 int main(int argc, char *argv[]) {
14     signal(SIGUSR1, handler1);
15     if ((pid = fork()) == 0) {
16         signal(SIGUSR1, handler2);
17         kill(getppid(), SIGUSR1);
18         while (true) {}
19     }
20     if (waitpid(-1, NULL, 0) > 0) {
21         counter += 1000;
22         printf("counter = %d\n", counter);
23     }
24     return 0;
25 }
```

- Now further assume the call to **exit(0)** has also been removed from the **handler2** function. Are there any other potential program outputs? If not, explain why. If so, what are they?
 - No other potential outputs, because:
 - ...
 - **counter = 1001** is always printed last, after the child process exits. It's possible that the child existed at the time the parent signaled it to inspire **handler2** to print a **10**, but that would happen before the **1001** is printed.

Plan For Today

- More Practice: Race Conditions
- **Break:** Announcements
- Threads



Announcements

- Assignment 4 released later today - Stanford Shell
- Midterm Next Friday
 - Please notify us of any OAE accommodations ASAP
 - More midterm logistics on Wednesday



Assignment 4: Stanford Shell

- Assignment 4 is a comprehensive test of your abilities to **fork** / **execvp** child processes and manage them through the use of signal handlers. It also tests your ability to use pipes.
- You will be writing a shell (demo: **assign4/samples/stsh_soln**)
 - The shell will keep a list of all background processes, and it will have some standard shell abilities:
 - you can quit the shell (using **quit** or **exit**)
 - you can bring them to the front (using **fg**)
 - you can continue a background job (using **bg**)
 - you can kill a set of processes in a pipeline (using **slay**) (this will entail learning about process groups)
 - you can stop a process (using **halt**)
 - you can continue a process (using **cont**)
 - you can get a list of jobs (using **jobs**)
 - You are responsible for creating pipelines that enable you to send output between programs, e.g.,
 - **ls | grep stsh | cut -d- -f2**
 - **sort < stsh.cc | wc > stsh-wc.txt**
 - You will also be handing off terminal control to foreground processes, which is new

Assignment 4: Stanford Shell

- Assignment 4 contains a lot of moving parts!
- Read through all the header files!
- You will only need to modify **stsh.cc**
- You can test your shell programmatically with **samples/stsh-driver**
- One of the more difficult parts of the assignment is making sure you are keeping track of all the processes you've launched correctly. This involves careful use of a **SIGCHLD** handler.
 - You will also need to use a handler to capture **SIGTSTP** and **SIGINT** to capture ctrl-Z and ctrl-C, respectively (notice that these don't affect your regular shell -- they shouldn't affect your shell, either).
- Another tricky part of the assignment is with the piping between processes. It takes time to understand what we are requiring you to accomplish
- There is a very good list of milestones in the assignment -- try to accomplish regular milestones, and you should stay on track.

Plan For Today

- More Practice: Race Conditions
- **Break:** Announcements
- Threads



Threads

A **thread** is an independent execution sequence within a single process.

- Most common: assign each thread to execute a single function in parallel
- Each thread operates within the same process, so they *share global data* (!) (text, data, and heap segments)
- They each have their own stack (e.g. for calls within a single thread)
- Execution alternates between threads as it does for processes
- Many similarities between threads and processes; in fact, threads are often called **lightweight processes**.



Threads vs. Processes

Processes:

- isolate virtual address spaces (good: security and stability, bad: harder to share info)
- can run external programs easily (fork-exec) (good)
- harder to coordinate multiple tasks within the same program (bad)

Threads:

- share virtual address space (bad: security and stability, good: easier to share info)
- can't run external programs easily (bad)
- easier to coordinate multiple tasks within the same program (good)



Threads

```
1 static void greeting() {
2     cout << oslock << "Hello, world!" << endl << osunlock;
3 }
4
5 static const size_t kNumFriends = 6;
6 int main(int argc, char *argv[]) {
7     cout << "Let's hear from " << kNumFriends << " threads." << endl;
8
9     thread friends[kNumFriends]; // declare array of empty thread handles
10
11     // Spawn threads
12     for (size_t i = 0; i < kNumFriends; i++) {
13         friends[i] = thread(greeting);
14     }
15
16     // Wait for threads
17     for (size_t i = 0; i < kNumFriends; i++) {
18         friends[i].join();
19     }
20
21     cout << "Everyone's said hello!" << endl;
22     return 0;
23 }
24
25
```



WARNING: Thread Safety and Standard I/O

- `operator<<`, unlike `printf`, isn't thread-safe.
 - Jerry Cain has constructed custom stream manipulators called `oslock` and `osunlock` that can be used to acquire and release exclusive access to an `ostream`.
 - These manipulators—which we can use by `#include`-ing "`ostreamlock.h`"—can be used to ensure at most one thread has permission to write into a stream at any one time.

Thread-Level Parallelism

- Threads allow a process to parallelize a problem across multiple cores
- Consider a scenario where we want to process 250 images and have 10 cores
- Completion time is determined by the slowest thread, so we want them to have equal work
 - Static partitioning: just give each thread 25 of the images to process. Problem: what if some images take much longer than others?
 - Work queue: have each thread fetch the next unprocessed image
- Here's our first stab at a **main** function.

```
int main(int argc, const char *argv[]) {  
    thread processors[10];  
    size_t remainingImages = 250;  
    for (size_t i = 0; i < 10; i++)  
        processors[i] = thread(process, 101 + i, ref(remainingImages));  
    for (thread& proc: processors) proc.join();  
    cout << "Images done!" << endl;  
    return 0;  
}
```

Thread Function

- The **processor** thread routine accepts an id number (used for logging purposes) and a reference to the **remainingImages**.
- It continually checks **remainingImages** to see if any images remain, and if so, processes the image and sends a message to **cout**
- **processImage** execution time depends on the image.
- Note how we can declare a function that takes a **size_t** and a **size_t&** as arguments

```
static void process(size_t id, size_t& remainingImages) {  
    while (remainingImages > 0) {  
        processImage(remainingImages);  
        remainingImages--;  
        cout << oslock << "Thread#" << id << " processed an image (" << remainingImages  
            << " remain)." << endl << osunlock;  
    }  
    cout << oslock << "Thread#" << id << " sees no remaining images and exits."  
        << endl << osunlock;  
}
```

- Discuss with your neighbor -- what's wrong with this code?

Race Condition

- Presented below right is the abbreviated output of a **imagethreads** run.
- In its current state, the program suffers from a serious race condition.
- Why? Because **remainingImages > 0** test and **remainingImages--** aren't atomic
- If a thread evaluates **remainingImages > 0** to be **true** and commits to processing an image, the image may have been claimed by another thread.
- This is a concurrency problem!
- Solution? Make the test and decrement *atomic* with a *critical section*
- Atomicity: externally, the code has either executed or not; external observers do not see any intermediate states mid-execution

```
myth60 ~../cs110/cthreads -> ./imagethreads
Thread# 109 processed an image, 249 remain
Thread# 102 processed an image, 248 remain
Thread# 101 processed an image, 247 remain
Thread# 104 processed an image, 246 remain
Thread# 108 processed an image, 245 remain
Thread# 106 processed an image, 244 remain
// 241 lines removed for brevity
Thread# 110 processed an image, 3 remain
Thread# 103 processed an image, 2 remain
Thread# 105 processed an image, 1 remain
Thread# 108 processed an image, 0 remain
Thread# 105 processed an image, 18446744073709551615 remain
Thread# 109 processed an image, 18446744073709551614 remain
```

Why Test and Decrement Is REALLY NOT Thread-Safe

- C++ statements aren't inherently atomic. Virtually all C++ statements—even ones as simple as **remainingImages--**—compile to multiple assembly code instructions.
- Assembly code instructions are atomic, but C++ statements are not.
- **g++** on the myths compiles **remainingImages--** to five assembly code instructions, as with:

```
0x0000000000401a9b <+36>:    mov     -0x20(%rbp), %rax
0x0000000000401a9f <+40>:    mov     (%rax), %eax
0x0000000000401aa1 <+42>:    lea     -0x1(%rax), %edx
0x0000000000401aa4 <+45>:    mov     -0x20(%rbp), %rax
0x0000000000401aa8 <+49>:    mov     %edx, (%rax)
```

- The first two lines drill through the **remainingImages** reference to load a copy of the **remainingImages** held on **main**'s stack. The third line decrements that copy, and the last two write the decremented copy back to the **remainingImages** variable held on **main**'s stack.
- The ALU operates on registers, but registers are private to a core, so the variable needs to be loaded from and stored to memory.
 - Each thread makes a local copy of the variable before operating on it
 - What if multiple threads all load the variable at the same time: they all think there's only 128 images remaining and process 128 at the same time

Mutual Exclusion

- A mutex is a type used to enforce *mutual exclusion*, i.e., a critical section
- Mutexes are often called locks
 - To be very precise, mutexes are one kind of lock, there are others (read/write locks, reentrant locks, etc.), but we can just call them locks in this course, usually "lock" means "mutex"
- When a thread locks a mutex
 - If the lock is unlocked the thread takes the lock and continues execution
 - If the lock is locked, the thread blocks and waits until the lock is unlocked
 - If multiple threads are waiting for a lock they all wait until lock is unlocked, one receives lock
- When a thread unlocks a mutex
 - It continues normally; one waiting thread (if any) takes the lock and is scheduled to run
- This is a subset of the C++ mutex abstraction: nicely simple!

```
class mutex {  
public:  
    mutex();           // constructs the mutex to be in an unlocked state  
    void lock();       // acquires the lock on the mutex, blocking until it's unlocked  
    void unlock();     // releases the lock and wakes up another threads trying to lock it  
};
```

Building a Critical Section with a Mutex

- **main** instantiates a mutex, which it passes (by reference!) to invocations of **process**.
- The **process** code uses this lock to protect **remainingImages**.
- Note we need to unlock on line 5 -- in complex code forgetting this is an easy bug

```
1 static void process(size_t id, size_t& remainingImages, mutex& counterLock) {
2     while (true) {
3         counterLock.lock();
4         if (remainingImages == 0) {
5             counterLock.unlock();
6             break;
7         }
8         processImage(remainingImages);
9         remainingImages--;
10        cout << oslock << "Thread#" << id << " processed an image (" << remainingImages
11            << " remain)." << endl << osunlock;
12        counterLock.unlock();
13    }
14    cout << oslock << "Thread#" << id << " sees no remaining images and exits."
15        << endl << osunlock;
16 }
17
18 int main(int argc, const char *argv[]) {
19     size_t remainingImages = 250;
20     mutex counterLock;
21     thread processors[10];
22     for (size_t i = 0; i < 10; i++)
23         agents[i] = thread(process, 101 + i, ref(remainingImages), ref(counterLock));
24     for (thread& agent: agents) agent.join();
25     cout << "Done processing images!" << endl;
26     return 0;
27 }
```

Critical Sections Can Be a Bottleneck

- The way we've set it up, only one thread agent can process an image at a time!
 - Image processing is actually serialized
- We can do better: serialize deciding which image to process and parallelize the actual processing
- Keep your critical sections as small as possible!

```
1 static void process(size_t id, size_t& remainingImages, mutex& counterLock) {
2     while (true) {
3         size_t myImage;
4
5         counterLock.lock();    // Start of critical section
6         if (remainingImages == 0) {
7             counterLock.unlock(); // Rather keep it here, easier to check
8             break;
9         } else {
10            myImage = remainingImages;
11            remainingImages--;
12            counterLock.unlock(); // end of critical section
13
14            processImage(myImage);
15            cout << oslock << "Thread#" << id << " processed an image (" << remainingImages
16            << " remain)." << endl << osunlock;
17        }
18    }
19    cout << oslock << "Thread#" << id << " sees no remaining images and exits."
20    << endl << osunlock;
21 }
```

Problems That Might Arise

- What if **processImage** can return an error?
 - E.g., what if we need to distinguish allocating an image and processing it
 - A thread can grab the image by decrementing **remainingImages** but if it fails there's no way for another thread to retry
 - Because these are threads, if one thread has a SEGV the whole process will fail
 - A more complex approach might be to maintain an actual queue of images and allow threads (in a critical section) to push things back into the queue
- What if image processing times are **highly** variable (e.g, one image takes 100x as long as the others)?
 - Might scan images to estimate execution time and try more intelligent scheduling
- What if there's a bug in your code, such that sometimes **processImage** randomly enters an infinite loop?
 - Need a way to reissue an image to an idle thread
 - An infinite loop of course shouldn't occur, but when we get to networks sometimes execution time can vary by 100x for reasons outside our control

Some Types of Mutexes

- Standard **mutex**: what we've seen
 - If a thread holding the lock tries to re-lock it, deadlock
- **recursive_mutex**
 - A thread can lock the mutex multiple times, and needs to unlock it the same number of times to release it to other threads
- **timed_mutex**
 - A thread can **try_lock_for** / **try_lock_until**: if time elapses, don't take lock
 - Deadlocks if same thread tries to lock multiple times, like standard mutex
- In this class, we'll focus on just regular **mutex**

How Do Mutexes Work?

- Something we've seen a few times is that you can't read and write a variable atomically
 - But a mutex does so! If the lock is unlocked, lock it
- How does this work with caches?
 - Each core has its own cache
 - Writes are typically write-back (write to higher cache level when line is evicted), not write-through (always write to main memory) for performance
 - Caches are *coherent* -- if one core writes to a cache line that is also in another core's cache, the other core's cache line is invalidated: this can become a performance problem
- Hardware provides atomic memory operations, such as compare and swap
 - `cas old, new, addr`
 - If `addr == old`, set `addr` to `new`
 - Use this as a single bit to see if the lock is held and if not, take it
 - If the lock is held already, then enqueue yourself (in a thread safe way) and tell kernel to sleep you
 - When a node unlocks, it clears the bit and wakes up a thread

**Questions about threads, mutexes,
race conditions, or critical sections?**