

# Lecture 15: Networking, Building an API

Principles of Computer Systems  
Spring 2019  
Stanford University  
Computer Science Department  
Instructors: Chris Gregg and  
Philip Levis



[PDF of this presentation](#)

# Lecture 15: API Servers, Threads, Processes

- Let's implement an API server that's architecturally in line with the way Google, Twitter, Facebook, and LinkedIn architect their own API servers.
- This example is inspired by a website called [Lexical Word Finder](#).
  - Our implementation assumes we have a standard Unix executable called **scrabbleword-finder**. The source code for this executable—completely unaware it'll be used in a larger networked application—can be found [right here](#).
  - **scrabble-word-finder** is implemented using only CS106B techniques—standard file I/O and procedural recursion with simple pruning.
  - Here are two abbreviated sample runs:

```
cgregg@myth61:$ ./scrabble-word-finder lexical
ace
// many lines omitted for brevity
lei
lex
lexica
lexical
li
lice
lie
lilac
xi
cgregg@myth61:$
```

```
cgregg@myth61:$ ./scrabble-word-finder network
en
// many lines omitted for brevity
wonk
wont
wore
work
worn
wort
wot
wren
wrote
cgregg@myth61:$
```

# Lecture 15: API Servers, Threads, Processes

- I want to implement an API service using HTTP to replicate what **scrabble-wordfinder** is capable of.
  - We'll expect the API call to come in the form of a URL, and we'll expect that URL to include the rack of letters.
  - Assuming our API server is running on **myth54:13133**, we expect **http://myth54:13133/lexical** and **http://myth54:13133/network** to generate the following payloads, in **JSON** format:

```
{
  "time": 0.223399,
  "cached": false,
  "possibilities": [
    "ace",
    // several words omitted
    "lei",
    "lex",
    "lexica",
    "lexical",
    "li",
    "lice",
    "lie",
    "lilac",
    "xi"
  ]
}
```

```
{
  "time": 0.242551,
  "cached": false,
  "possibilities": [
    "en",
    // several words omitted
    "wonk",
    "wont",
    "wore",
    "work",
    "worn",
    "wort",
    "wot",
    "wren",
    "wrote"
  ]
}
```

# Lecture 15: API Servers, Threads, Processes

- One might think to cannibalize the code within `scrabble-word-finder.cc` to build the core of `scrabble-word-finder-server.cc`.
- Reimplementing from scratch is wasteful, time-consuming, and unnecessary. `scrabble-word-finder` already outputs the primary content we need for our payload. We're packaging the payload as JSON instead of plain text, but we can still tap `scrabble-word-finder` to generate the collection of formable words.
- Can we implement a server that leverages existing functionality? Of course we can!
- We can just leverage our `subprocess_t` type and `subprocess` function from Assignment 3.

```
struct subprocess_t {
    pid_t pid;
    int supplyfd;
    int ingestfd;
};

subprocess_t subprocess(char *argv[],
    bool supplyChildInput, bool ingestChildOutput) throw (SubprocessException);
```

# Lecture 15: API Servers, Threads, Processes

- Here is the core of the `main` function implementing our server:

```
int main(int argc, char *argv[]) {
    unsigned short port = extractPort(argv[1]);
    int server = createServerSocket(port);
    cout << "Server listening on port " << port << "." << endl;
    ThreadPool pool(16);
    map<string, vector<string>> cache;
    mutex cacheLock;
    while (true) {
        struct sockaddr_in address;
        // used to surface IP address of client
        socklen_t size = sizeof(address); // also used to surface client IP address
        bzero(&address, size);
        int client = accept(server, (struct sockaddr *) &address, &size);
        char str[INET_ADDRSTRLEN];
        cout << "Received a connection request from "
             << inet_ntop(AF_INET, &address.sin_addr, str, INET_ADDRSTRLEN) << "." << endl;
        pool.schedule([client, &cache, &cacheLock] {
            publishScrabbleWords(client, cache, cacheLock);
        });
    }
    return 0;
}
```

# Lecture 15: API Servers, Threads, Processes

- The second and third arguments to **accept** are used to surface the IP address of the client.
- Ignore the details around how I use **address**, **size**, and the **inet\_ntop** function until the next lecture, when we'll talk more about them. Right now, it's a neat-to-see!
- Each request is handled by a dedicated worker thread within a **ThreadPool** of size 16.
- The thread routine called **publishScrabbleWords** will rely on our **subprocess** function to marshal plain text output of scrabble-word-finder into JSON and publish that JSON as the payload of the HTTP response.
- The next slide includes the full implementation of **publishScrabbleWords** and some of its helper functions.
- Most of the complexity comes around the fact that I've *elected* to maintain a cache of previously processed letter racks.

# Lecture 15: API Servers, Threads, Processes

- Here is `publishScrabbleWords`:

```
static void publishScrabbleWords(int client, map<string, vector<string>>& cache,
                                mutex& cacheLock) {
    sockbuf sb(client);
    iosockstream ss(&sb);
    string letters = getLetters(ss);
    sort(letters.begin(), letters.end());
    skipHeaders(ss);
    struct timeval start;
    gettimeofday(&start, NULL); // start the clock
    cacheLock.lock();
    auto found = cache.find(letters);
    cacheLock.unlock(); // release lock immediately, iterator won't be invalidated by competing find calls
    bool cached = found != cache.end();
    vector<string> formableWords;
    if (cached) {
        formableWords = found->second;
    } else {
        const char *command[] = {"/scrabble-word-finder", letters.c_str(), NULL};
        subprocess_t sp = subprocess(const_cast<char **>(command), false, true);
        pullFormableWords(formableWords, sp.ingestfd);
        waitpid(sp.pid, NULL, 0);
        lock_guard<mutex> lg(cacheLock);
        cache[letters] = formableWords;
    }
    struct timeval end, duration;
    gettimeofday(&end, NULL); // stop the clock, server-computation of formableWords is complete
    timersub(&end, &start, &duration);
    double time = duration.tv_sec + duration.tv_usec/1000000.0;
    ostringstream payload;
    constructPayload(formableWords, cached, time, payload);
    sendResponse(ss, payload.str());
}
```

# Lecture 15: API Servers, Threads, Processes

- Here's the `pullFormableWords` and `sendResponse` helper functions.

```
static void pullFormableWords(vector<string>& formableWords, int ingestfd) {
    stdio_filebuf<char> inbuf(ingestfd, ios::in);
    istream is(&inbuf);
    while (true) {
        string word;
        getline(is, word);
        if (is.fail()) break;
        formableWords.push_back(word);
    }
}

static void sendResponse(iosockstream& ss, const string& payload) {
    ss << "HTTP/1.1 200 OK\r\n";
    ss << "Content-Type: application/javascript; charset=UTF-8\r\n";
    ss << "Content-Length: " << payload.size() << "\r\n";
    ss << "\r\n";
    ss << payload << flush;
}
```



# Lecture 15: API Servers, Threads, Processes

## HTTP status ranges in a nutshell:

1xx: hold on

2xx: here you go

3xx: go away

4xx: you f up

5xx: I f up

-via @abt\_programming

# Lecture 15: API Servers, Threads, Processes

- Finally, here are the `getLetters` and the `constructPayload` helper functions. I omit the implementation of `skipHeaders`—you saw it with `web-get`—and `constructJSONArray`, which you're welcome to view [right here](#).

```
static string getLetters(iosockstream& ss) {
    string method, path, protocol;
    ss >> method >> path >> protocol;
    string rest;
    getline(ss, rest);
    size_t pos = path.rfind("/");
    return pos == string::npos ? path : path.substr(pos + 1);
}

static void constructPayload(const vector<string>& formableWords, bool cached, double time,
                             ostream& payload) {
    payload << "{" << endl;
    payload << "  \"time\": " << time << ", " << endl;
    payload << "  \"cached\": " << boolalpha << cached << ", " << endl;
    payload << "  \"possibilities\": " << constructJSONArray(formableWords, 2) << endl;
    payload << "}" << endl;
}
```

- Our `scrabble-word-finder-server` provided a single API call that resembles the types of API calls afforded by Google, Twitter, or Facebook to access search, tweet, or friend-graph data.

# Lecture 15: API Servers, Threads, Processes

- It turns out we actually wrote a web-ready program. Let's look at some HTML and Javascript:

```
1 <!DOCTYPE html>
2 <html>
3   <head>
4     <title>Scrabble Word Finder</title>
5   </head>
6   <body>
7     Letters: <input type="text" id="letters" name="letters"><br>
8     <input type="submit" value="Submit" onclick="getWords()">
9     <p>Time:</p><span id="words_time"></span>
10    <p>Scrabble words:</p>
11    <div id="scrabble_words"></div>
12
13    <script>
14      function getWords(){
15        let letters = document.getElementById("letters").value;
16        let scrabble_words = fetch("http://myth59.stanford.edu:13133/"+letters,{method:"GET"})
17          .then(data=>{return data.json()})
18          .then(res=>{
19            console.log(res);
20            document.getElementById("words_time").innerText = res['time']+"sec";
21            possibilitiesStr = "";
22            for (var i=0; i < res.possibilities.length; i++) {
23              possibilitiesStr += res.possibilities[i]+"<br>";
24            }
25            document.getElementById("scrabble_words").innerHTML = possibilitiesStr;
26          })
27          .catch(error=>console.log(error))
28      }
29    </script>
30  </body>
31 </html>
```

# Lecture 15: API Servers, Threads, Processes

- We can put the html file in a Stanford web location (e.g., <http://web.stanford.edu/class/cs110/scrabble-word-finder.html>)
- As long as we are running our scrabble server on myth59, port 13133, we can get words:

Letters: croaatb

Submit

Time:

0.041713sec

Scrabble words:

aa  
aba  
abator  
abo  
abort  
acrobat  
act  
acta  
actor  
aorta  
ar  
arb  
arc  
arco  
art  
at  
ba  
baa

Letters: croaatb

Submit

Time:

2.3e-05sec

Scrabble words:

aa  
aba  
abator  
abo  
abort  
acrobat  
act  
acta  
actor  
aorta  
ar  
arb  
arc  
arco  
art  
at  
ba  
baa

- We have a legitimate back-end server. It has a cache, just like built it to have.
- Are many servers written in C++?
- Surprisingly, yes -- it is fast, which often makes it better than Python, or Node, or PHP.

# Lecture 15: Network System Calls, Library Functions

- Hostname Resolution: IPv4
- Linux C includes directives to convert host names (e.g. "**www.facebook.com**") to IPv4 address (e.g. "**31.13.75.17**") and vice versa. Functions called **gethostbyname** and **gethostbyaddr**, while technically deprecated, are still so prevalent that you should know how to use them.
- In fact, your B&O textbook only mentions these deprecated functions:

```
struct hostent *gethostbyname(const char *name);  
struct hostent *gethostbyaddr(const char *addr, int len, int type);
```

- Each function populates a statically allocated **struct hostent** describing some host machine on the Internet.
  - **gethostbyname** assumes its argument is a host name (e.g. "**www.google.com**").
  - **gethostbyaddr** assumes the first argument is a binary representation of an IP address (e.g. not the string "**171.64.64.137**", but the base address of a character array with ASCII values of 171, 64, 64, and 137 laid down side by side in **network byte order**. For IPv4, the second argument is usually 4 (or rather, **sizeof(struct in\_addr)**) and the third is typically the **AF\_INET** constant.

# Lecture 15: Network System Calls, Library Functions

- Hostname Resolution: IPv4
  - The **struct hostent** record packages all of the information about a particular host:

```
struct in_addr {
    unsigned int s_addr // four bytes, stored in network byte order (big endian)
};
struct hostent {
    char *h_name;
    // official name of host
    char **h_aliases;
    // NULL-terminated list of aliases
    int h_addrtype;
    // host address type (typically AF_INET for IPv4)
    int h_length;
    // address length (typically 4, or sizeof(struct in_addr) for IPv4)
    char **h_addr_list; // NULL-terminated list of IP addresses
}; // h_addr_list is really a struct in_addr ** when hostent contains IPv4 addresses
```

- The **struct in\_addr** is a one-field record modeling an IPv4 address.
  - The **s\_addr** field packs each figure of a dotted quad (e.g. 171.64.64.136) into one of its four bytes. Each of these four numbers can range from 0 up through 255.
- The **struct hostent** is used for all IP addresses, not just IPv4 addresses. For non-IPv4 addresses, **h\_addrtype**, **h\_length**, and **h\_addr\_list** carry different types of data than they do for IPv4

# Lecture 15: Network System Calls, Library Functions

Users prefer the host naming scheme behind "**www.facebook.com**", but network communication ultimately works with IP addresses like "31.13.75.17".

- Not surprisingly, **gethostbyname** and **gethostbyaddr** are used to manage translations between the two.
- Here's the core of larger program (full program [here](#)) that continuously polls the users for hostnames and responds by publishing the set of one or more IP addresses each hostname is bound to:

```
static void publishIPAddressInfo(const string& host) {
    struct hostent *he = gethostbyname(host.c_str());
    if (he == NULL) { // NULL return value means resolution attempt failed
        cout << host << " could not be resolved to an address. Did you mistype it?" << endl;
        return;
    }

    cout << "Official name is \"" << he->h_name << "\"" << endl;
    cout << "IP Addresses: " << endl;
    struct in_addr **addressList = (struct in_addr **) he->h_addr_list;
    while (*addressList != NULL) {
        char str[INET_ADDRSTRLEN];
        cout << "+ " << inet_ntop(AF_INET, *addressList, str, INET_ADDRSTRLEN) << endl;
        addressList++;
    }
}
```

# Lecture 15: Network System Calls, Library Functions

## Hostname Resolution: IPv4

`h_addr_list` is typed to be a `char *` array, implying it's an array of C strings, perhaps dotted quad IP addresses. However, that's not correct. For IPv4 records, `h_addr_list` is an array of `struct in_addr *s`.

The `inet_ntop` function places a traditional C string presentation of an IP address into the provided character buffer, and returns the base address of that buffer.

The while loop crawls over the `h_addr_list` array until it lands on a `NULL`.

```
static void publishIPAddressInfo(const string& host) {
    struct hostent *he = gethostbyname(host.c_str());
    if (he == NULL) { // NULL return value means resolution attempt failed
        cout << host << " could not be resolved to an address. Did you mistype it?" << endl;
        return;
    }

    cout << "Official name is \"" << he->h_name << "\"" << endl;
    cout << "IP Addresses: " << endl;
    struct in_addr **addressList = (struct in_addr **) he->h_addr_list;
    while (*addressList != NULL) {
        char str[INET_ADDRSTRLEN];
        cout << "+ " << inet_ntop(AF_INET, *addressList, str, INET_ADDRSTRLEN) << endl;
        addressList++;
    }
}
```



# Lecture 15: Network System Calls, Library Functions

## Hostname Resolution: IPv4

- A sample run of our hostname resolver is presented on the right.
- In general, you see that most of the hostnames we recognize are in fact the officially recorded hostnames.
- **www.yale.edu**, **www.facebook.com**, and **www.wikipedia.org** are exceptions. It looks like Yale relies on a content delivery network called Cloudflare, and **www.yale.edu** is catalogued as an alias.
- Google's IP address is different by geographical location, which is why it exposes only one IP address.

```
myth61$ ./resolve-hostname
Welcome to the IP address resolver!
Enter a host name: www.google.com
Official name is "www.google.com"
IP Addresses:
+ 216.58.192.4
Enter a host name: www.coinbase.com
Official name is "www.coinbase.com"
IP Addresses:
+ 104.16.9.251
+ 104.16.8.251
Enter a host name: www.yale.edu
Official name is "www.yale.edu.cdn.cloudflare.net"
IP Addresses:
+ 104.16.140.133
+ 104.16.141.133
Enter a host name: www.facebook.com
Official name is "star-mini.c10r.facebook.com"
IP Addresses:
+ 31.13.70.36
Enter a host name: www.wikipedia.org
Official name is "dyna.wikimedia.org"
IP Addresses:
+ 198.35.26.96
Enter a host name:
All done!
myth61$
```

# Lecture 15: Network System Calls, Library Functions

## Hostname Resolution: IPv6

- Because IPv4 addresses are 32 bits, there are  $2^{32}$ , or roughly 4 billion different IP addresses. That may sound like a lot, but it was recognized decades ago that we'd soon run out of IPv4 addresses.
- In contrast, there are 340,282,366,920,938,463,374,607,431,768,211,456 IPv6 addresses. That's because IPv6 addresses are 128 bits.
- Here are a few IPv6 addresses:
  - Google's 2607:f8b0:4005:80a::2004
  - MIT's 2600:1406:1a:396::255e and 2600:1406:1a:38d::255e
  - Berkeley's 2600:1f14:436:7801:15f8:d879:9a03:eec0 and 2600:1f14:436:7800:4598:b474:29c4:6bc0
  - The White House's 2600:1406:1a:39e::fc4 and 2600:1406:1a:39b::fc4

A more generic version of **gethostbyname**—inventively named **gethostbyname2**—can be used to extract IPv6 address information about a hostname.

```
struct hostent *gethostbyname2(const char *name, int af);
```

# Lecture 15: Network System Calls, Library Functions

## Hostname Resolution: IPv6

- There are only two valid address types that can be passed as the second argument to `gethostbyname2`: `AF_INET` and `AF_INET6`.
  - A call to `gethostbyname2(host, AF_INET)` is equivalent to a call to `gethostbyname(host)`
  - A call to `gethostbyname2(host, AF_INET6)` still returns a `struct hostent *`, but the struct `hostent` is populated with different values and types:
    - the `h_addrtype` field is set to `AF_INET6`,
    - the `h_length` field houses a 16 (or rather, `sizeof(struct in6_addr)`), and
    - the `h_addr_list` field is really an array of `struct in6_addr` pointers, where each `struct in6_addr` looks like this:

```
struct in6_addr {  
    u_int8_t s6_addr[16]; // 16 bytes (128 bits), stored in network byte order  
};
```

# Lecture 15: Network System Calls, Library Functions

## Hostname Resolution: IPv6

- Here is the `IPv6` version of the `publishIPAddressInfo` we wrote earlier (we call it `publishIPv6AddressInfo`).

```
static void publishIPv6AddressInfo(const string& host) {
    struct hostent *he = gethostbyname2(host.c_str(), AF_INET6);
    if (he == NULL) { // NULL return value means resolution attempt failed
        cout << host << " could not be resolved to an address. Did you mistype it?" << endl;
        return;
    }

    cout << "Official name is \"" << he->h_name << "\"" << endl;
    cout << "IPv6 Addresses: " << endl;
    struct in6_addr **addressList = (struct in6_addr **) he->h_addr_list;
    while (*addressList != NULL) {
        char str[INET6_ADDRSTRLEN];
        cout << "+ " << inet_ntop(AF_INET6, *addressList, str, INET6_ADDRSTRLEN) << endl;
        addressList++;
    }
}
```

- Notice the call to `gethostbyname2`, and notice the explicit use of `AF_INET6`, `struct in6_addr`, and `INET6_ADDRSTRLEN`.
- Full program is [right here](#).

# Lecture 15: Network System Calls, Library Functions

## Hostname Resolution: IPv6

- A sample run of our IPv6 hostname resolver is presented below.
  - Note that many hosts aren't IPv6-compliant yet, so they don't admit IPv6 addresses.

```
myth61$ ./resolve-hostname6
Welcome to the IPv6 address resolver!
Enter a host name: www.facebook.com
Official name is "star-mini.c10r.facebook.com"
IPv6 Addresses:
+ 2a03:2880:f131:83:face:b00c:0:25de
Enter a host name: www.microsoft.com
Official name is "e13678.dspb.akamaiedge.net"
IPv6 Addresses:
+ 2600:1406:1a:386::356e
+ 2600:1406:1a:397::356e
Enter a host name: www.google.com
Official name is "www.google.com"
IPv6 Addresses:
+ 2607:f8b0:4005:801::2004
Enter a host name: www.berkeley.edu
Official name is "www-production-1113102805.us-west-2.elb.amazonaws.com"
IPv6 Addresses:
+ 2600:1f14:436:7800:4598:b474:29c4:6bc0
+ 2600:1f14:436:7801:15f8:d879:9a03:eec0
Enter a host name: www.stanford.edu
www.stanford.edu could not be resolved to an address. Did you mistype it?
Enter a host name:
All done!
myth61$
```